

# Minimum Description Length Principle

## With An Application For Credit Risk Models

Paul van Leeuwen ([paul.vanleeuwen@devolksbank.nl](mailto:paul.vanleeuwen@devolksbank.nl))

7 February 2023

## Introduction

## What is MDL?

## MDL in Practice

## Approaches compared

## Discussion

## Conclusion

## Appendix

# Introduction

# The Speaker

- Paul van Leeuwen, employee of KRT RDS 5.0 as of September 2022.
- Up to 2017 part of the Modelling team at dVB.
- After that as Model Validator at Achmea and Lead Data Scientist at Wageningen University & Research.
- Now self-employed and currently working in the financial sector and providing R workshops.

## Why this subject?

- Improve current approaches and modelling techniques.
- Stay in the forefront of statistical innovation.
- Connection with other realms of statistics.

# What is MDL?

# Introduction

- The Minimum Description Length (MDL) principle aims to describe the data and its description mechanism with the smallest possible information 'length'.
- Applications (among else):
  - model selection (e.g. what order of the Markov model family do we want?);
  - deal with overfitting (e.g. how many explanatory variables to include);
  - exploratory data analysis (what prior knowledge can we confirm?).
- Close ties with frequentist statistics, Bayesian statistics, and machine learning.
- Why is MDL relatively unknown?
  - MDL is the intersection of advanced measure theory, information theory, and statistics.
  - For a decent introduction into MDL, see (Grünwald 2007).

# Example of dealing with overfitting


What polynomial generated by this dataset?






# Example of dealing with overfitting

linear fit  $\hat{y} = \beta_0 + \beta_1 x$




# Example of dealing with overfitting

quadratic fit  $\hat{y} = \beta_0 + \beta_1 x + \beta_2 x^2$



# Example of dealing with overfitting

10th order polynomial fit  $\hat{y} = \beta_0 + \beta_1 x + \dots + \beta_{10} x^{10}$




# Example of dealing with overfitting

data generated by  $x^3 + 2x^2 + 5 + \epsilon$ ,  $\epsilon \sim N(0, 1)$

linear

quadratic

10th order polynomial



## How does MDL work?

- Patterns or regularities in the data can be described with less information 'length' compared to the data alone.
- less information 'length' = compression
- Choose the model that gives the shortest description of the data.
- Note that MDL is an approach, not an algorithm.
  - The modeller has to make choices to implement the MDL principle.

## Example of MDL

- Consider three data-generating processes (dgp's) that generate each a binary sequence of length 1000:
  - 100010001000100010001000100010001000100 ... 1000
  - 1001111110111100011001110001010100110001101 ... 1010
  - 0100010000001011100100000000001101000110000 ... 0100

## Example of MDL

- Consider three data-generating processes (dgp's) that generate each a binary sequence of length 1000:
  1. 100010001000100010001000100010001000100 ... 1000
  2. 10011111101111100011001110001010100110001101 ... 1010
  3. 0100010000001011100100000000001101000110000 ... 0100
- All sequences written out take 1000 bits to be reproduced.
- However, because of regularities present in the dgp's, we require less bits to reproduce the same sequences.
- Question: how many bits of Matlab code does each sequence take to be *exactly* reproduced?

## Example of MDL

- The sequences as before:

1. 100010001000100010001000100010001000100010001000100 ... 1000:



$$1 - \frac{184}{1000} = 81.6\%!$$

## Example of MDL

- The sequences as before:
  1. 10001000100010001000100010001000100010001000100 ... 1000:  
 repetition of [1 0 0 0] 250 times.  
 Matlab-code:  

```
textDgp_1 = 'repmat([1 0 0 0],1,250)'
```

```
whos_textDgp_1 = whos('textDgp_1')
```

```
whos_textDgp_1.bytes
```

 yields 23 bytes = 184 bits, a compression ratio of  
 $1 - \frac{184}{1000} = 81.6\%$ !
  2. 1001111110111100011001110001010100110001101 ... 1010:

## Example of MDL

- The sequences as before:

1. 100010001000100010001000100010001000100010001000100 ... 1000:  
repetition of [1 0 0 0] 250 times.

Matlab-code:

```
textDgp_1 = 'repmat([1 0 0 0],1,250)'
```

```
whos_textDgp_1 = whos('textDgp_1')
```

```
whos_textDgp_1.bytes
```

yields 23 bytes = 184 bits, a compression ratio of

$$1 - \frac{184}{1000} = 81.6\%$$

2. 1001111110111100011001110001010100110001101 ... 1010:  
a coin toss with heads (0) or tails (1); no compression possible  
because of the randomness involved.
3. 0100010000001011100100000000001101000110000 ... 0100:

- The sequences as before:
  1. 1000100010001000100010001000100010001000100 ... 1000:  
repetition of [1 0 0 0] 250 times.  
Matlab-code:  

```
textDgp_1 = ' repmat([1 0 0 0],1,250) '
whos_textDgp_1 = whos('textDgp_1')
whos_textDgp_1.bytes
```

 yields 23 bytes = 184 bits, a compression ratio of  
 $1 - \frac{184}{1000} = 81.6\%$ !
  2. 1001111110111100011001110001010100110001101 ... 1010:  
a coin toss with heads (0) or tails (1); no compression possible  
because of the randomness involved.
  3. 0100010000001011100100000000001101000110000 ... 0100:  
a roll with a four-sided die with outcomes 2, 3, 4 assigned to 0  
and outcome 1 assigned to 1; no compression possible because  
of the randomness involved.

## The Principle of MDL

- The more randomness involved, the less data compression is possible.
- When the restriction of exact reproduction is alleviated we may obtain some data compression.
- Dgp 2 (the coin toss) implies no data compression.
  - As the data sequence is completely random.
- Dgp 3 (the 4-sided die) implies some possible data compression.
  - Although exact reproduction is not possible the set of this type of data sequences requires around 821 bits.
- Most datasets are almost incompressible.
  - Only a small fraction can be significantly compressed.

## Description Methods

- A description method in conjunction with a coding scheme maps an object to a number of bits.
  - The coding alphabet  $\mathbb{B}$  can be binary ( $\mathbb{B} = \{0, 1\}$ ), the Western alphabet ( $\mathbb{B} = \{a, b, \dots, z\}$ ), etc.
  - The mapping is a one-many relation and many-one.
  - Mathematically: a dataset  $D = (x_1, \dots, x_n)$  with  $x_i \in \mathbb{B}$  from a sample space  $\mathcal{X}^n$  is mapped to  $\mathbb{R}_+$  by  $L: \mathcal{X}^n \mapsto \mathbb{R}_+$ .
- Any alphabet is fine.
- For example, to map an integer from  $\{1, \dots, n\}$  in a uniform way, we need  $\log_2 n$  bits.
  - Note that  $n$  has to be known in advance.
  - For example, take  $n = 64$ . Then we have 64 binary data sequences of length  $\log_2(64) = 6$ .

## The Principle of MDL (continued)

- To describe any dataset we need  $L(D, H)$  bits.
  - Both the description method  $H$  and the data  $D$  require storage space.
- The MDL principle employed for model selection is to minimise the sum of
  - the number of bits to encode the description mechanism  $L(H)$  and
  - the number of bits to encode, with the description mechanism  $H$ , the data observed  $L(D|H)$ .
- Information 'length' is this sum  $L(H) + L(D|H)$ .
- The former is the *model complexity*, the latter the *fit of the data*.
- The MDL principle is to choose the model specification as to minimise this sum.

## The Principle of MDL (continued)

- More mathematically, given a set of candidate models  $\mathcal{H} = \{\mathcal{H}_1, \mathcal{H}_2, \dots\}$ , select the optimal model  $\mathcal{H}^* \in \mathcal{H}$  as

$$\mathcal{H}^* := \arg \min_{H \in \mathcal{H}} L(D, H) = \arg \min_{H \in \mathcal{H}} \{L(H) + L(D|H)\}$$


- Note the resemblance with penalised model fitting, such as LASSO:
  - with LASSO we apply cross-validation to minimise

$$\arg \min_{\beta} \left\{ \frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1 \right\}$$


with the data part  $\frac{1}{n} \|\mathbf{y} - \mathbf{X}\beta\|_2^2$  ( $L(D|H)$  in MDL) and  $\lambda \|\beta\|_1$  the model complexity part ( $L(H)$  in MDL).



# Example of LASSO



## Example of dealing with overfitting (continued)



## Questions to be answered

- How does MDL work in practice?
- What is the best approach: frequentist statistics, Bayesian statistics, machine learning, or MDL?
- The coding system to put down the description mechanism should not matter. How can we choose a universal programming language?
  - For example, whether we use Matlab, R or whatever programming language should not matter for the number of bits at use.
- What are good encoding systems?
- What patterns are generalisable and what not?
- How to incorporate prior knowledge?

# MDL in Practice

## Dow Jones Industrial Average

- How can we employ MDL to compress a dataset?
  - Example taken from (Hansen and Yu 2001).
- Take the log daily return  $R_t$  and the volatility  $V_t$  of the Dow Jones Industrial Average (DJIA).
  - $t$  runs from  $t_0 = \text{July 1962}$  until  $T = \text{June 1988}$ , i.e. 6,430 trading days.
  - $R_t = P_t - P_{t-1}$  with  $P_t$  the logarithm of the DJIA at day  $t$ .
  - $V_t = 0.9V_{t-1} + 0.1R_t^2$  and  $V_0$  the variance of the series  $P_t$ .
- $R_t$  has a corresponding indicator: 1 (0) when  $R_t > R_{t-1}$  ( $R_t < R_{t-1}$ ).
  - Analogously for  $V_t$ .
  - Two binary strings of length  $6,430 - 1 = 6,429$ .
  - $R_t$  has 3,181 (49.49%) ups.
  - $V_t$  has 2,023 or (31.47%) ups.

# Dow Jones Industrial Average

logarithm of the closing of the Dow Jones Industrial Average



## Dow Jones Industrial Average

- Without data compression we require 6,429 bits per series.
- Using MDL we save 10% on the volatility series  $V_t$  while gaining nothing on  $P_t$ .
- $n$  is known in advance so costs us  $\lceil \log_2 n \rceil$  bits.
- Model the up or down indicator as a Bernoulli probability distribution with probability  $p$  on success.
  - Maximum likelihood yields  $\hat{p} = k/n$  with  $k$  the number of successes.
- It takes us  $-\log_2(k/n)$  bits to model a success and  $-\log_2(1 - k/n)$  bits to model a failure.
- In total we need, ignoring rounding errors,

$$\log_2(n) + k \left[ -\log_2 \left( \frac{k}{n} \right) \right] + (n - k) \left[ -\log_2 \left( 1 - \frac{k}{n} \right) \right]$$

## Dow Jones Industrial Average

- With  $n = 6,429$  and
  - for  $R_t$  we have  $k = 3181$  and  $\hat{p} = \frac{k}{n} = \frac{3181}{6429} = 0.49$  we need
$$\log_2(6429) + 3181[-\log_2 0.49] + 3248[-\log_2 0.51] = 6442 \text{ bits}$$
  - for  $V_t$  we have  $k = 2023$  and  $\hat{p} = \frac{k}{n} = \frac{2023}{6429} = 0.31$  we need
$$\log_2(6429) + 2023[-\log_2 0.31] + 4406[-\log_2 0.69] = 5789 \text{ bits}$$
- In conclusion, for  $R_t$  we have gained nothing, even lost some bits, but the improvement for  $V_t$  is substantial.
  - The savings in storage increases as well as  $n$  increases or  $k/n$  approaches 0 or 1.
- In general, data compression leads to good model performance.
  - Note that data compression is not the main goal.



# Approaches compared

## How do the approaches compare: Frequentist

- Also known as orthodox or non-Bayesian.
- Main goal: models should be consistent.
  - At least asymptotically: as  $n \rightarrow \infty$  the estimated model coefficients should converge to the 'true' model parameters.
  - In the example above, one would expect  $\hat{\beta} = (5 \ 0 \ 2 \ 1)$  for  $n \rightarrow \infty$ .
- When the model assumptions are fulfilled, this approach yields the best possible predictions.
- Pros: widely known and applicable, easy calculation, models are good with infinite amount of data.
- Cons: assumptions are not realistic (e.g. almost no residual follows a normal probability distribution).

## How do the approaches compare: Bayesian

- Main goal: incorporate prior knowledge by assigning a prior probability distribution.
  - Once the prior probability distribution is assigned, the posterior probability distribution can be derived / calculated.
  - Frequentist models consider the data random and the parameters fixed.
  - Bayesian models consider the data and the parameters random.
  - In practice, asymptotically frequentist and Bayesian models converge.
- Pros: mathematically elegant and precise, encompasses the frequentist approach.
- Cons: hard to come up with a useful prior probability distribution.

# How do the approaches compare: Machine Learning

- Main goal: find and fit a model that is able to generalise from train to test data.
  - No assumptions about the data generating process.
- Pros: no assumptions required, good model performance.
- Cons: hard to explain predictions, hard to study the data generating process.

## How do the approaches compare: MDL

- Main goal: infer useful information from the data and use that to achieve good data compression.
  - Good data compression implies good learning.
  - No assumptions about the data generating process required but are allowed.
- Pros: no assumptions required, has best of all worlds (Bayesian and Machine Learning): employ prior knowledge when viable, good statistical properties, and decent performance.
- Cons: relatively new in literature, intersection of state-of-the-art knowledge of information theory, measure theory and statistics.

Introduction  
ooo

What is MDL?  
oooooooooooooooooooo

MDL in Practice  
ooooo

Approaches compared  
ooooo

Discussion  
●o

Conclusion  
oooo

Appendix  
oooo

# Discussion

## Application to Expected Credit Loss

- Suppose we apply the MDL principle on the calculation of the Expected Credit Loss.
  - No coding scheme discussed here, just the principle.
- Discussion:

*take a moving average of the realised credit losses to calculate the Expected Credit Loss*
- Pros:
  - Occam's Razor: simpler is better.
  - Saves us a complete IFRS team to perform tedious calculations.
- Cons:
  - No new developments (e.g. Corona) taken into consideration. Can we with our current models?
  - No intermediate analyses available (e.g. use PD to rank customers in need).

Introduction  
○○○

What is MDL?  
○○○○○○○○○○○○○○○○○○○○

MDL in Practice  
○○○○○

Approaches compared  
○○○○○

Discussion  
○○

**Conclusion**  
●○○○

Appendix  
○○○○

# Conclusion



## Summary

- MDL is a guidance for model selection, an instrument against overfitting and exploratory data analysis.
- Sound statistics and decent model performance could make MDL a challenger for other statistical approaches.
- Much more to be explored.

## What's next?

- A workshop to take a deepdive into the applications of MDL, backed up by statistical theory.
- Take a look at the open questions raised in this presentation.
  - The coding system to put down the description mechanism should not matter. How can we choose a universal programming language?
  - What are good encoding systems?
  - What patterns are generalisable and what not?
  - How to incorporate prior knowledge?
- Dedicated application to IFRS and Regulatory Capital calculations.
- Please let Eelko ([eelko.ubels@devolksbank.nl](mailto:eelko.ubels@devolksbank.nl)) or me ([paul.vanleeuwen@devolksbank.nl](mailto:paul.vanleeuwen@devolksbank.nl)) know whether you would like to join!

Introduction  
○○○

What is MDL?  
○○○○○○○○○○○○○○○○○○○○

MDL in Practice  
○○○○○

Approaches compared  
○○○○○

Discussion  
○○

**Conclusion**  
○○○●

Appendix  
○○○○

# Questions

?

# Appendix

## Kolmogorov complexity

- Any application of MDL should not be affected by the programming language of choice.
- Andrey Kolmogorov postulated the invariance theorem:  
*given any description of an object in a description language  $A$ ,  
said description may be used in the optimal description  
language  $B$  with a constant overhead*
- The description follows two steps:
  - describe optimal language  $B$  in language  $A$ ;
  - describe object in optimal language  $B$  via language  $A$ .
- Comparable with a user-friendly programming language  $A$  (e.g. Matlab, R, Python) and whose code is compiled in an efficient programming language  $B$  (e.g. C++, C, FORTRAN).
- Part 1. is overhead independent of the object to be described, hence as the object size grows the overhead becomes small.

## Kolmogorov complexity (example)

- Suppose we would like to describe the binary data sequence 1 0 0 0 in  $C$  via Matlab. How many bits does it take?
- First translate any code from  $C$  to Matlab needed to describe 1 0 0 0.
  - E.g.  $C$  has a certain mapping procedure from the data type bit to its memory.
- The corresponding number of bits of the Matlab-code is  $L_{C \mapsto M}$ , e.g. 10k bits.
- To describe  $D = 1\ 0\ 0\ 0$  in  $C$  via Matlab takes  $L_{C \mapsto M} + L_M(D)$  bits.
- But  $L_{C \mapsto M}$  is independent of the object to be described.
  - Whether we put in the works of Shakespeare or Hello, World,  $L_{C \mapsto M}$  remains the same.

## Bibliography

Grünwald, Peter D. 2007. *The Minimum Description Length Principle*. MIT press.

Hansen, Mark H, and Bin Yu. 2001. "Model Selection and the Principle of Minimum Description Length." *Journal of the American Statistical Association* 96 (454): 746–74.